



A four-seat dark factory that built Pocketful

Team **Revived** · WeAreDevelopers × BAND Dark Factory · track pocketful

4 / 4

stages claimed


4

human messages


2 h 18 min

\$35.72 at list price

Submitted run, room pocketful-final, 5 Oct 2026. Our own harness runs and ccusage, not the judges' scores.



Most of each later stage's tests are hidden



"A green run on the shipped checks is not evidence of a stage."

Participant guide, "Do not write to the tests"

When one model writes both the code and its checks, its blind spots land in both.

So we split them: an independent seat writes hidden checks from the spec alone, behind a wall the builder cannot cross.

Four Claude Code seats in one Band room

Human: one task per stage, nothing else → sent by `dispatch_stages.py` after the coordinator's final report

 **coordinator**
claude-sonnet-5
Hands out work, reviews the diff, rules disputes from spec quotes, closes the stage.

 **builder**
claude-opus-5-5
Builds each stage whole: API first, then markup with the designer's class names.

 **designer**
claude-opus-5-5
Design system, stylesheet, SVG art, mockups; checks layout in Chromium at 375 and 1280 px.

 **tester**
claude-opus-5-5
Interface contract first, then holdout scenarios from the spec alone; runs the gate.

 **holdout/ — readable by the tester only**
`permissions.deny: Read(../../holdout/**) for coordinator, builder and designer. Enforced by config, not by prompt; tested with Read, cat and Python, in bypass mode.`

```
Gate output, builder-view format: no test code crosses the wall

BUILD: ok
HOLDOUT: 573/574 requirements green
  FAILED "A stage-4 service must accept
exports produced by the same team's stages
1-3, retaining settlement membership,
corrections and snapshots."
SHIPPED CHECKS: ok
VERDICT: RED
```

The builder only ever learns *which spec sentence* it fails. In the submitted run this RED went from the tester to the coordinator for a ruling (#1370 → #1386). Every view forwarded to the builder was GREEN.

- 1

Coordinator copies the stage forward, one task per seat
- 2

Tester commits `contract.json` in its own short turn
- 3

Builder gets the full spec + contract; designer publishes `system.md`
- 4

Every builder commit: gate + layout check
- 5

Fix, or dispute with a quote; ruling in `decisions.md`
- 6

Close: gate green + LAYOUT: clean, ≤ 5 review fixes

ⓘ HOW IT CATCHES BAD WORK

Five checks — and the two things they still missed



Independent holdout

Written from the spec alone; other seats denied by config.



Quotes checked by script

coverage.py rejects any quote not in the spec and lists uncovered sentences.



One gate

gate.py: Docker build, scenarios, shipped checks, one verdict.



Failures are signals

A builder may dispute; the coordinator rules from a quote.



Layout by browser

layout_check.py: overflow, overlap, clipped text, <40 px controls.

✓ Caught in the submitted run



Blank screen on cold load · stage 2

Designer's review on 2ef0c94: layout clean, but the page stayed blank until /me answered. Builder added a boot state **79 s** later (0229cfa).



Unguarded type assertion · stage 4

Coordinator's diff review: parseBatchItem could panic (500) instead of a clean 422. Fixed in 30bfb5a, re-gated green.



The tester's own mistakes · stages 2–3

Three red rounds were scenario bugs (a route the spec doesn't pin, a sign flipped, a same-second race). The tester fixed its tests; correct code was not changed.

ⓘ Missed — kept visible on purpose



A wrong ruling · stage 4

The holdout caught that snapshots did not survive importing a stage-3 export. The tester asked for a ruling; the coordinator ruled that a stage-3 export never carried snapshots, and the scenario was narrowed. The spec does require it — *“retaining settlement membership, corrections and snapshots”*. Our post-run hidden suite still fails that check (91/92). The arbiter is a model too.



Isolated mode · stage 3 sample

In the judges' isolated mode one stage-3 sample test fails (5/6, in folders 3 and 4): same-timestamp entries sort by payment id, not creation order. Seats ran checks in host mode, where the tie never happens. Fix for the factory: give seats `--mode isolated`.

The seats did the work together, and without us

4

human messages in the whole run — one task per stage, sent by a script. No steering, approvals, hints or reruns.

39 commits by the seats

coordinator		22
tester		8
builder		6
designer		3

Plus 1 setup commit. Each seat has its own git identity, so every line of stage-*/ traces to a seat and a room message.

Room log

1,531 events, 72 text messages: coordinator **37** · tester **20** · builder **7** · designer **4** · human **4**. All events by seat, same order: 695 · 397 · 292 · 143.

Handoffs carry the whole task; review changes the work

16:55:42	human → coordinator	Stage-1 task with the full spec (25,924 chars)
16:57:54	coordinator → tester	"Write and commit the contract, report it, nothing else."
16:59:10	tester → coordinator	Contract committed (590644f), 16 operations — 76 s later
16:59:38	coordinator → builder	The whole spec (26,002 chars) plus the contract path
17:32:23	designer → builder	Stage-2 design system and exact markup, before any screen exists
17:54:15	designer → builder	Review: blank screen on cold load → fix 0229cfa at 17:55:34
18:55:09	tester → coordinator	Gate RED 573/574, "I need a ruling" → ruling at 18:58:29, recorded in decisions.md
19:02:41	coordinator → builder	Maintainability review → fix 30bfb5a at 19:03:56, re-gated green



All four stages, 2 h 18 min end to end

Stage 1	Stage 2	Stage 3	Stage 4
17:00	17:30	18:00	18:30
19:00			
Stage 1 32 min Gategreen, round 2 Shipped checks147/147 Holdout233/233 · 252 tests Our hidden suite—	Stage 2 47 min Gategreen, round 3 Shipped checks147 · 35/35 Holdout420/420 · 374 tests Layoutclean, 3 reviews	Stage 3 27 min Gategreen, round 3 Shipped checks147 · 35 · 6/6* Holdout530/530 · 457 tests Our hidden suite149/149	Stage 4 30 min Gategreen, round 3 Shipped checks147 · 35 · 6* · 5/5 Holdout574/574 · 490 tests Our hidden suite91/92
4 / 4 stages claimed; none overshoots	100% spec coverage at close (688/688 spec units)	2 rulings recorded (1 tooling, 1 spec)	2 h 18 min 16:55 → 19:14 UTC; stages sum to 136 min

* Host mode. In the judges' isolated mode the stage-3 sample suite is 5/6 in folders 3 and 4; every folder still claims its stage. Shipped checks are not a blind score here: the seats ran them in the gate. Hidden suite = independent checks we wrote after the run.

6

Pocketful, as the factory shipped it



375 px
Nav becomes four tabs. No sideways scroll.

Money in tabular figures; held funds in indigo with a lock.

LAYOUT: clean
at 375 and 1280 px


Designer first. 40 mockups (20 screen states × 2 widths), the design system and the class names were committed before the builder wrote a screen.

Builder owns behaviour. Markup uses only the designer's classes; no visual CSS of its own. One stylesheet, app.css.


Checked by a browser. Every screen in headless Chromium at both widths: overflow, overlap, clipped text, 40 px controls, 14 px text.

Every state the stage-2 spec names looks different


“Available, held, pending, loading, successful, refused and uncertain states must be visually distinct” — by colour, icon and words.




● Loading
Boot state the designer's review asked for.




● Held
Indigo + lock; expired and voided muted.




● Refused
Says why, and what is available.



● Uncertain
“Retrying is safe and won't do it twice.”



● Empty
Custom illustration and the next action.



● Preview
Split shares before anything is sent.

We measured the alternatives before choosing

Submitted run · 4 stages

Spend, list price (ccusage)

\$35.72

Opus \$26.87 · builder, tester, designerSonnet \$8.86

Wall time

2 h 18 min

Stages 32 · 47 · 27 · 30 min

Development run, same design

Factory **\$42.76** vs one Opus session **\$10.43**: about 4× the cost, level on hidden checks (239/241 vs 238/241), and it fixed a UI defect the single session left in.

What we tried first


Development runs · pocketful stage 1 · shipped tests hidden from the factory, scored after · bar = shipped checks of 147

One Opus session 7 min · \$1.50 · stage 2: 35/35		147
One Haiku session 11 min · \$1.07 · stage 2: 1/35		137
3 Sonnet seats, template mandates 20 min · ~ \$11 · review missed a real bug		145
5 seats, 15 cards, Haiku builder > 25 min · \$7.78		stopped early
Haiku builder + 2 Haiku testers ≈ 55 min · \$11.45 · 3 rounds		57 → 132
Opus tester + 2 Haiku builders ≈ 45 min · LLM holdout picked the worse one		126 / 139

⚠ Cheap testers wrote status-only checks: their gate said **21/26** green while the service passed **57/147**. So: frontier builder, independent Opus tester, rules in config — prose rules broke with every model.

Same mandates, another problem

The mandates name no endpoint, field or error code. To check, we pointed them unchanged at the other track.



Five steps from FACTORY.md

- 1

Install **Band Desktop**, **Docker**, **Claude Code**, **Python 3.12+**; create a Claude token and store it in the keychain as `df-claude-oauth`.
- 2

`factory/setup/install.sh <workspace> <result-repo> --create` writes the seat launchers and the two isolated configs (holdout denied), and creates the four Band agents.
- 3

Copy **mandates/** and **factory/** into the result repository and commit.
- 4

Create a room and add the four seats. Start in a fresh usage window; detach old sessions (the Band daemon has a 256 open-file limit).
- 5

Per stage, fill `factory/TASK_TEMPLATE.md` and send it to the coordinator. **Send nothing else** until its final report.



Pocketful

Thanks — Revived

Mandates, tools, every stage's evidence and the full room log, unchanged.

github.com/Whom-m0rty/revived-pocketful-dark-factory

FACTORY.md	design, costs, what failed
mandates/	the four seats, generic
factory/tools/	gate, coverage, layout check
factory/stage-N/	contract, design, decisions, evidence
room.json	submitted run's room log